

The background image shows a busy logistics yard. On the left, a yellow container is being hoisted by a crane and loaded onto a blue truck. To the right, there are tall stacks of colorful shipping containers in shades of yellow, red, and blue. The sky is filled with soft, colorful clouds, suggesting a sunrise or sunset. The overall scene is industrial and dynamic.

# Cost Prediction for Logistic Company

INCLASS REGRESSION PROBLEM FOR AML-1413 FALL 2023

BY TEAM GAMMA B





# *TEAM MEMBERS*

---

AKASH(  
AKHIL  
JAGADEESH  
KAMESWARA  
KAPIL  
SATHWICK



## *STEPS INVOLVED*

---

- Data Validation & cleansing*
- Data preprocessing*
- Exploratory Data Analysis*

# Data Validation and Cleansing



# Data Validation & Cleansing

- Dimensions of the data
- The Train dataset (38999 Rows and 12 Columns)
- The Test dataset (802 Rows and 11 Columns)
- Converting date to date-time format
- Checking missing values and its percentage

## Data Validation and Cleaning

```
In [3]: # Let's Look The Dimensions Of The Data:
print(f'The Train dataset Contain {train.shape[0]} Rows and {train.shape[1]} Columns')
```

The Train dataset Contain 38999 Rows and 12 Columns

```
In [5]: # Let's Look The Dimensions Of The Data:
print(f'The Test dataset Contain {test.shape[0]} Rows and {test.shape[1]} Columns')
```

The Test dataset Contain 802 Rows and 11 Columns

```
In [6]: #Check Data Types
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38999 entries, 0 to 38998
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   trip                   38999 non-null  object
1   date                   38999 non-null  object
2   dayPart                38999 non-null  object
3   exWeatherTag          4882 non-null   object
4   originLocation         38999 non-null  object
5   destinationLocation    38999 non-null  object
6   distance               38999 non-null  int64
7   type                   3748 non-null   object
8   weight                38999 non-null  int64
9   packageType            2500 non-null   object
10  carrier                38999 non-null  object
11  cost                   38999 non-null  float64
dtypes: float64(1), int64(2), object(9)
```

# Converting date into datetime Format

```
## Converting date to datetime format and extracting the temporal features  
train['date'] = pd.to_datetime(train['date'], format='%Y-%m-%d')  
train['year'] = train['date'].dt.year  
train['month'] = train['date'].dt.month  
train['day'] = train['date'].dt.day  
  
test['date'] = pd.to_datetime(test['date'], format='%Y-%m-%d')  
test['year'] = test['date'].dt.year  
test['month'] = test['date'].dt.month  
test['day'] = test['date'].dt.day
```

# Check for missing values in the data

*## Checking missing values and visualizing it*

```
missing_counts = test.isnull().sum()
present_values=test.notnull().sum()
unique_value_counts = test.nunique()
total_count = len(test)
missing_percentage= (missing_counts / total_count) * 100
present_percentage= (present_values / total_count) * 100
unique_percentage= (unique_value_counts / total_count) * 100
```

*# Plotting the missing values % in the dataset*

```
plt.figure(figsize=(20, 10))
ax=sns.barplot(x=missing_counts.index, y=missing_counts.values, palette='viridis')
```

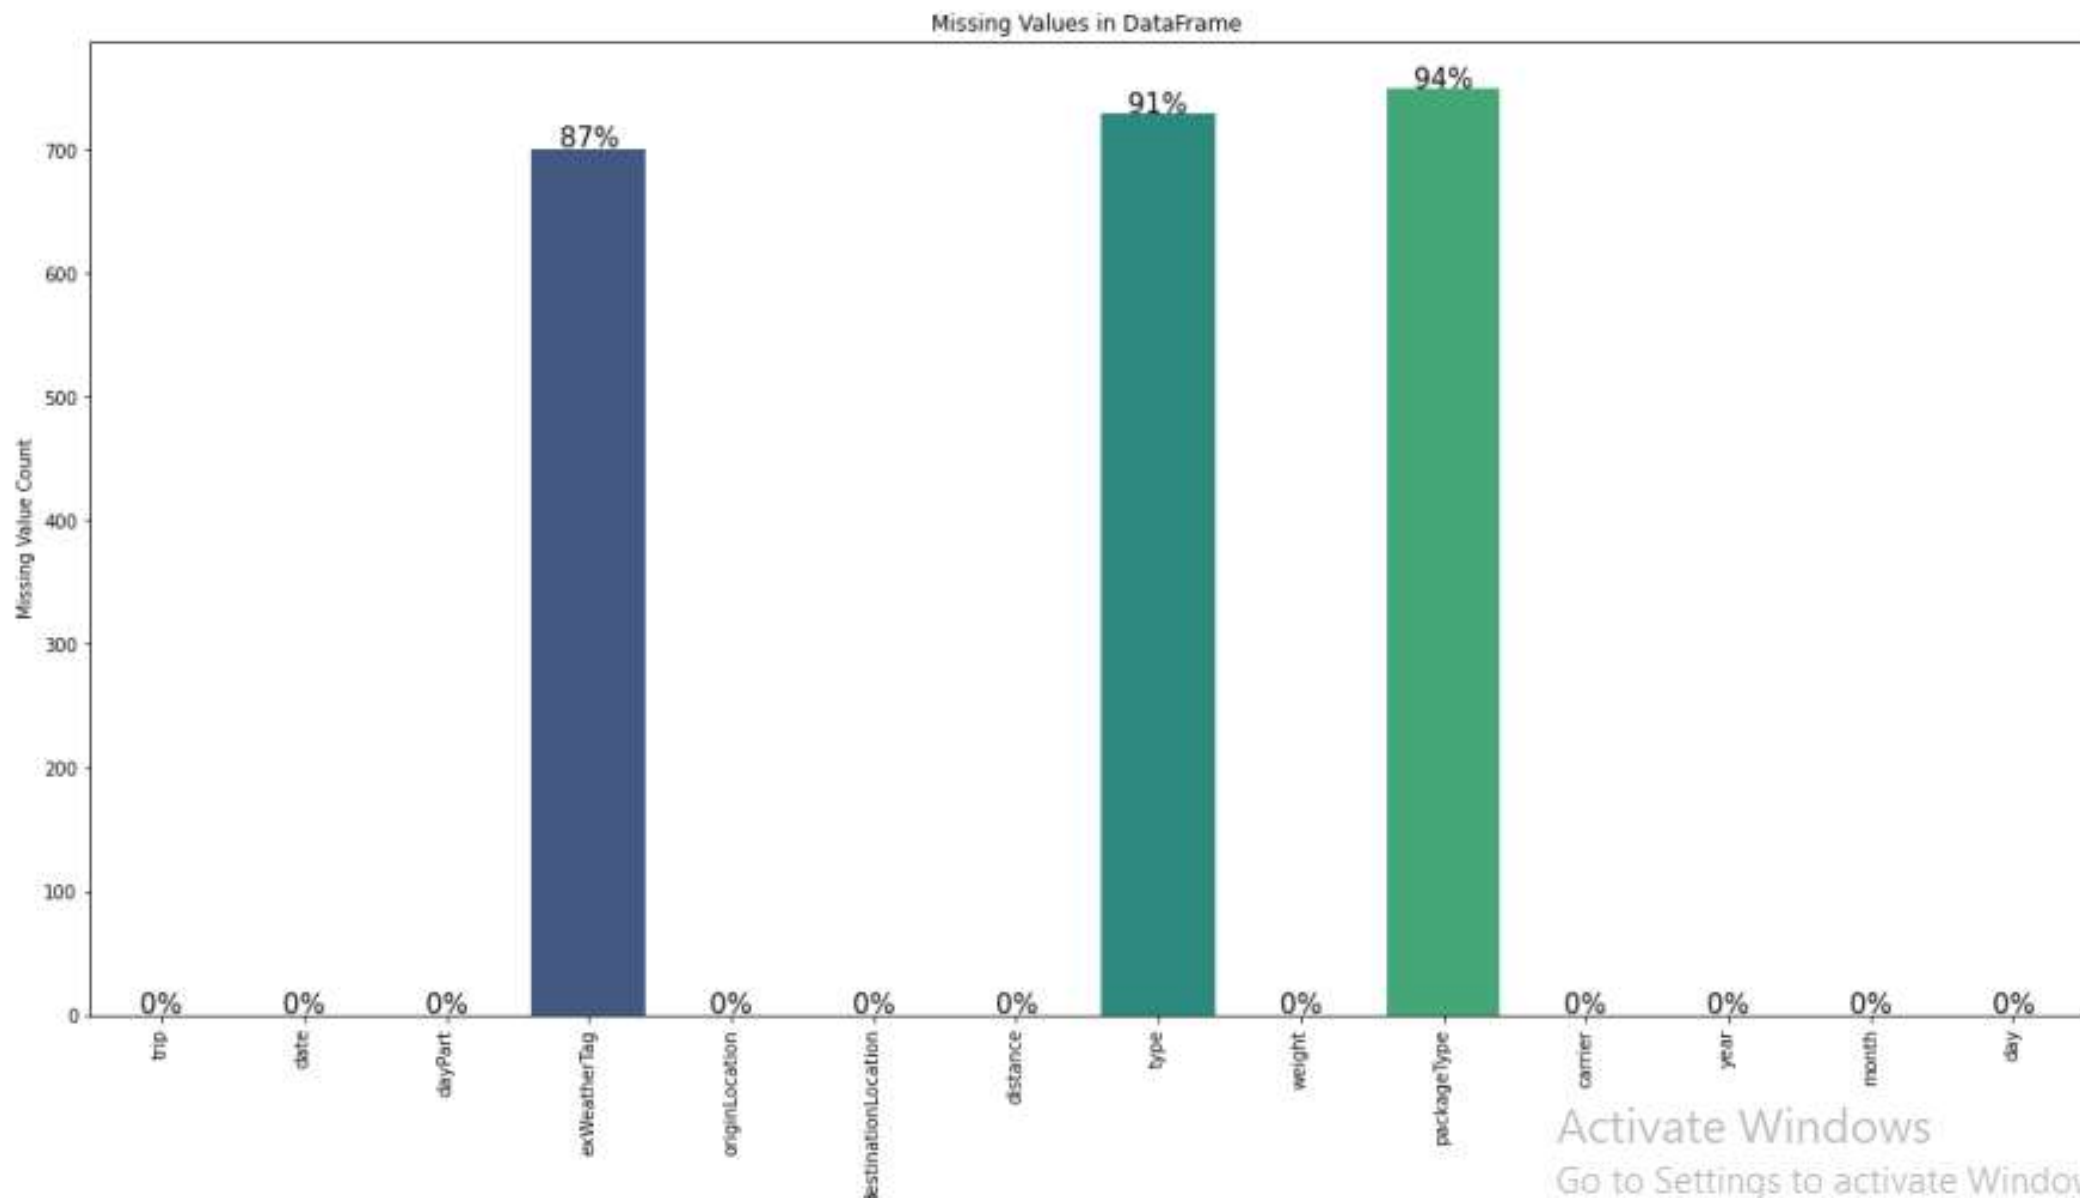
*# Add Labels and title*

```
plt.title('Missing Values in DataFrame')
plt.xlabel('Columns')
plt.ylabel('Missing Value Count')
plt.xticks(rotation=90)
```

```
for p, label in zip(ax.patches, missing_percentage):
```

```
    ax.annotate(f'{label:.0f}%', (p.get_x() + p.get_width() / 2., p.get_height()),
               ha='center', va='center', fontsize=15, color='black', xytext=(0, 5), textcoords='offsetpoints')
```

```
plt.show()
```



Activate Windows  
Go to Settings to activate Windows



# Imputing features having high percentage of missing values

- Type, packageType and exWeatherTag have high missing %, so let's impute the missing values as 'Not\_available' category since they seem to be important information.

In [10]:

```
# Sample of the unique elements of the features
Unique_values_list=[]
for column in train.columns:
    unique_values = train[column].unique().tolist()
    num_unique=train[column].nunique()
    if num_unique==1:
        Unique_values_list.append(column)
    print("{} Unique values in '{}': {} ".format(num_unique,column,unique_values[:5]))
print(Unique_values_list)
```

```
38999 Unique values in 'trip': ['t52712528', 't29859381', 't25702332', 't27713405', 't49439220']
1074 Unique values in 'date': [1504656000000000000, 1508544000000000000, 1500076800000000000, 1508630400000000000, 1576022400000000000]
2 Unique values in 'dayPart': ['night', 'day']
2 Unique values in 'exWeatherTag': [nan, 'snow', 'heat']
9 Unique values in 'originLocation': ['S4', 'S8', 'S9', 'S6', 'S7']
9 Unique values in 'destinationLocation': ['D7', 'D1', 'D5', 'D4', 'D2']
17 Unique values in 'distance': [2200, 1800, 2800, 3200, 2000]
1 Unique values in 'type': ['expedited', nan]
499 Unique values in 'weight': [50, 12, 1, 43, 3]
1 Unique values in 'packageType': [nan, 'TT']
4 Unique values in 'carrier': ['D', 'B', 'C', 'A']
3665 Unique values in 'cost': [68.41315193, 36.45064919, 9.05793866, 57.32008718, 77.2637771]
3 Unique values in 'year': [2017, 2019, 2018]
12 Unique values in 'month': [9, 10, 7, 12, 6]
30 Unique values in 'day': [6, 21, 15, 22, 11]
['type', 'packageType']
```

Type, packageType and exWeatherTag have high missing %, so let's impute the missing values as 'Not\_available' category since they seem to be important information.

## Feature Engineering

In [14]:

```
train = train.apply(lambda x: x.fillna('NA'))
test = test.apply(lambda x: x.fillna('NA'))
```

# Importing holidays data set

Since it is a time series dataset, let's add date related features such as holiday\_or\_not, weekday\_or\_not

```
In [19]: # using holidays library to get the holidays
import holidays

holiday_list = list()
print("--- Canada ---")
for date in holidays.Canada(years=[2014,2015,2016,2017,2018,2019], observed=True).items():
    print(str(date[0]), date[1])
    holiday_list.append([date[0], date[1]])
Holiday_dataset=pd.DataFrame(holiday_list,columns=['Date','Holiday'])

Holiday_dataset['Date'] = pd.to_datetime(Holiday_dataset['Date'])

# Create 'day_Holiday_or_not' column in train and test DataFrames
train['day_Holiday_or_not'] = train['date'].isin(Holiday_dataset['Date']).astype(int)
test['day_Holiday_or_not'] = test['date'].isin(Holiday_dataset['Date']).astype(int)
train['is_weekday'] = train['date'].apply(lambda x: 1 if x.weekday() < 5 else 0)
test['is_weekday'] = test['date'].apply(lambda x: 1 if x.weekday() < 5 else 0)
train['weekday']=train['date'].dt.day_name()
test['weekday']=test['date'].dt.day_name()
```

```
--- Canada ---
2016-01-01 New Year's Day
2016-03-25 Good Friday
2016-07-01 Canada Day
2016-09-05 Labour Day
2016-12-25 Christmas Day
2016-12-26 Christmas Day (Observed)
2017-01-01 New Year's Day
2017-01-02 New Year's Day (Observed)
2017-04-14 Good Friday
2017-07-01 Canada Day
2017-09-04 Labour Day
2017-12-25 Christmas Day
2018-01-01 New Year's Day
2018-03-30 Good Friday
2018-07-01 Canada Day
2018-09-03 Labour Day
2018-12-25 Christmas Day
```



The image features a server rack on the left side, with several units visible. The units have blue indicator lights and perforated metal doors. The background is a soft-focus bokeh of yellow and blue circles, suggesting a data center environment. The text "Data Preprocessing" is centered in the middle of the image in a white, italicized serif font.

# *Data Preprocessing*

# Data Preprocessing

- Splitting into train and test (X\_train, X\_test, Y\_train, Y\_test)
- Normalizing the variables using standard scalar

```
val_trip=test.pop('trip')
test.drop(columns=['date'],inplace=True)
train.drop(columns=['trip','date'],inplace=True)
```

```
Y=train.pop('cost')
X=train.copy()
```

*#Splitting train and test*

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3, random_state=42)
```

```
numerical_columns = X_train.select_dtypes(include=['int64', 'float64']).columns
```

*# Normalizing the variables*

```
pipeline = Pipeline([
    ('std_scalar', StandardScaler())
])
```

```
X_train[numerical_columns] = pipeline.fit_transform(X_train[numerical_columns])
```

```
X_test[numerical_columns] = pipeline.transform(X_test[numerical_columns])
```

```
test[numerical_columns]=pipeline.transform(test[numerical_columns])
```





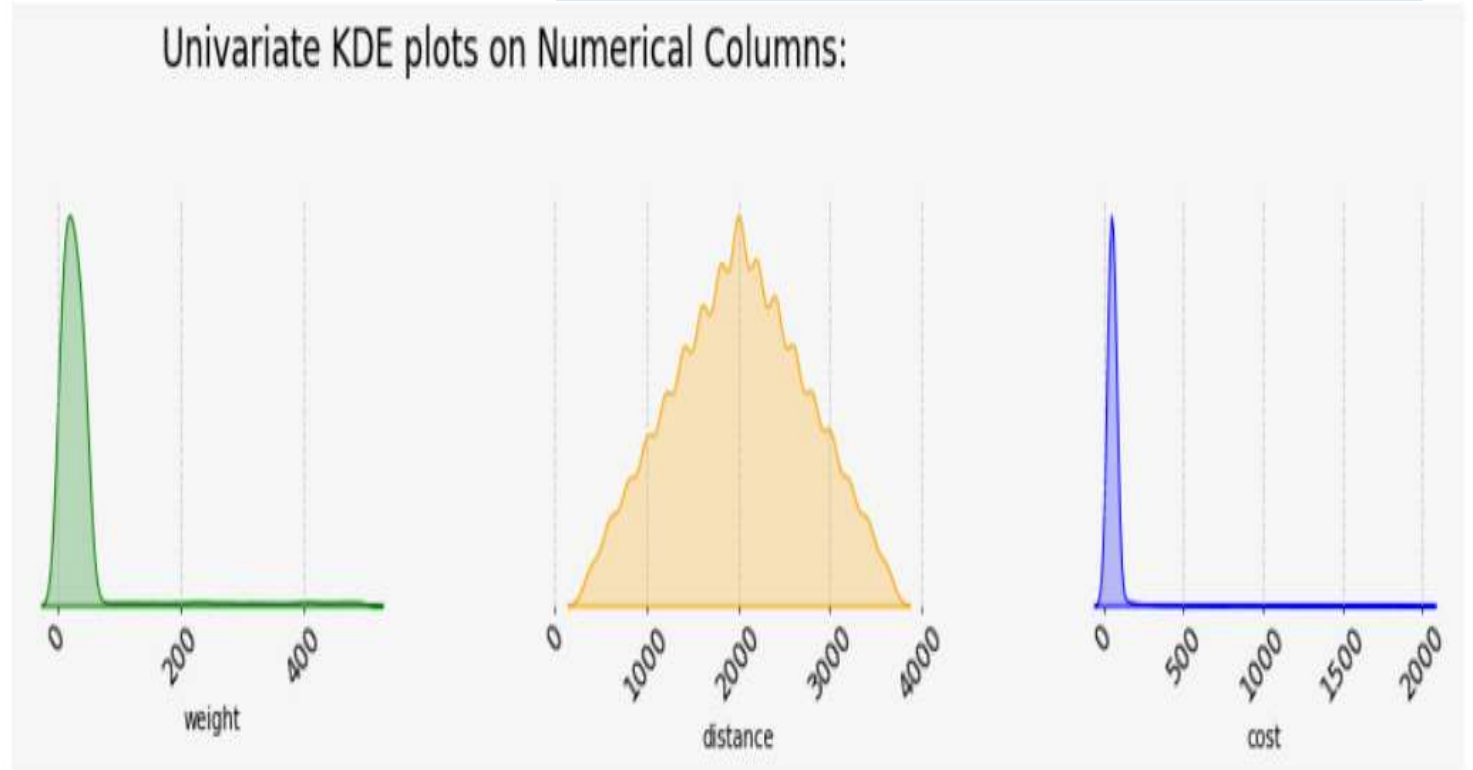
# EDA

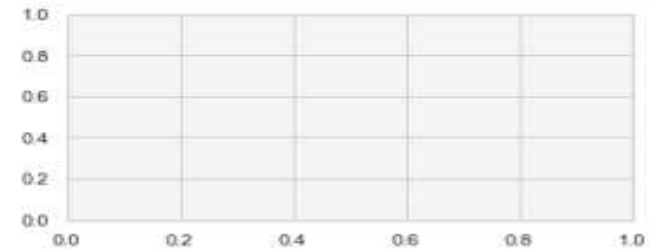
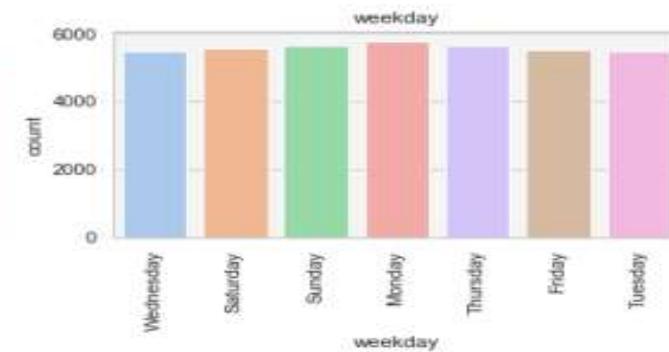
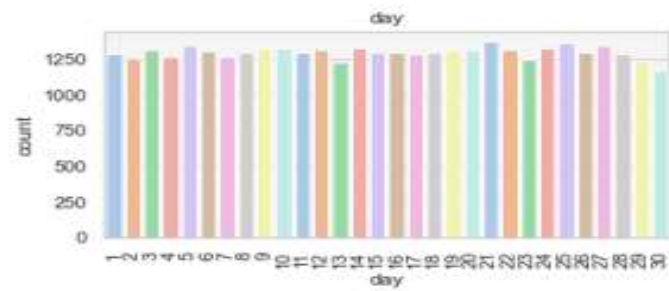
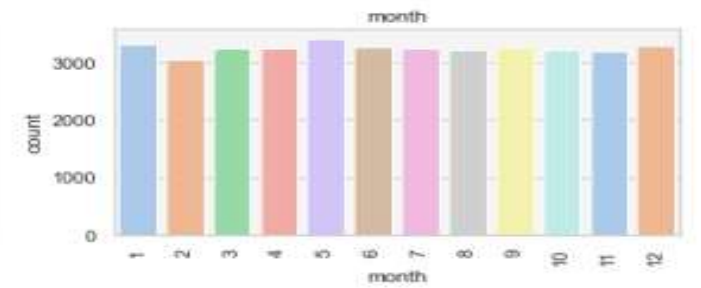
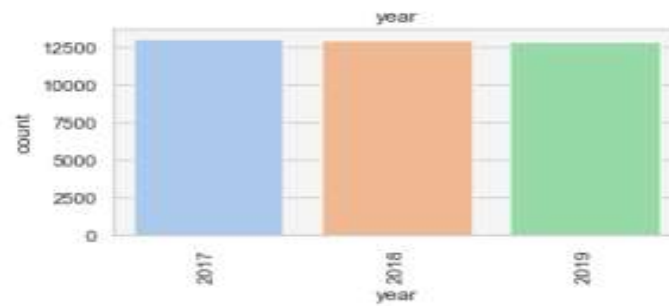
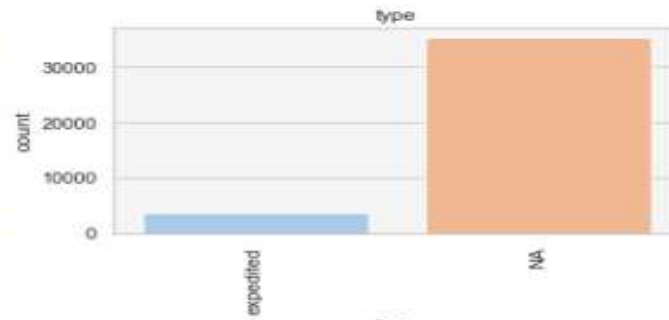
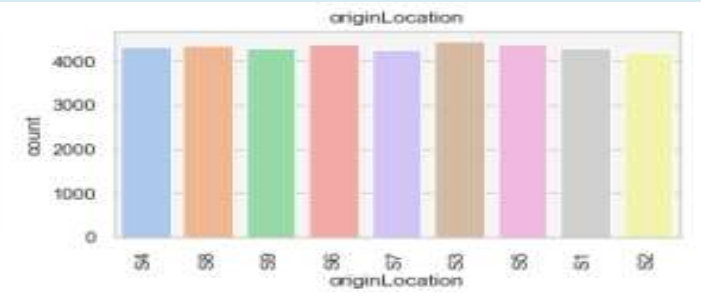
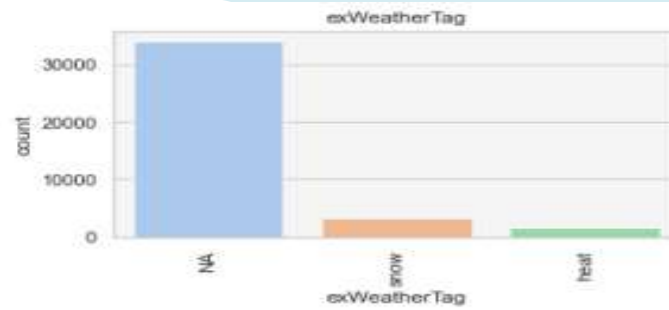
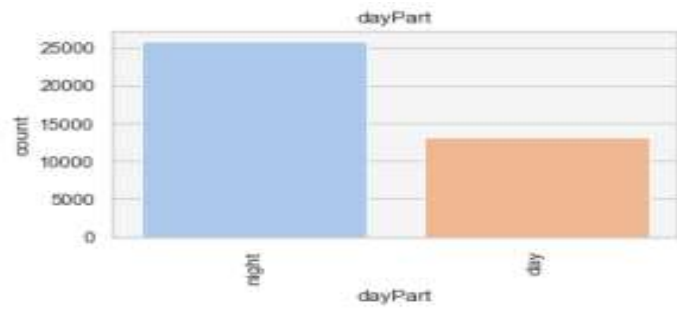
UNIVARIATE ANALYSIS  
BIVARIATE ANALYSIS



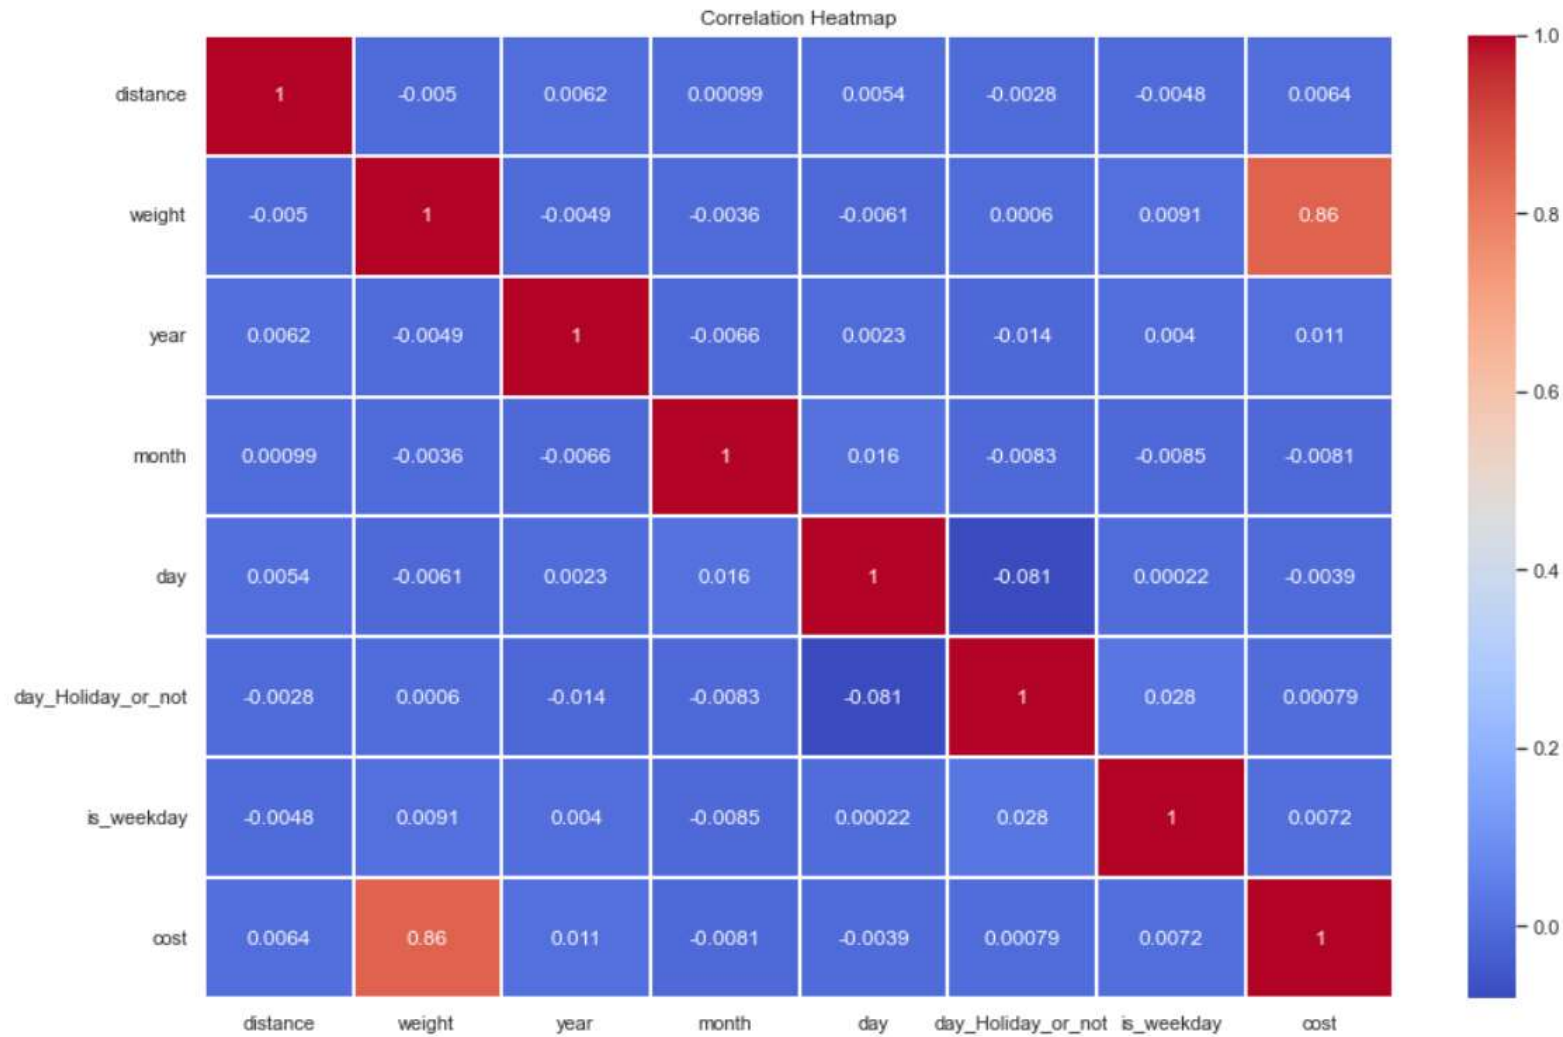
# Univariate Analysis

PLOTTING THE UNIVARIATE KDE  
PLOTS ON NUMERICAL COLUMNS





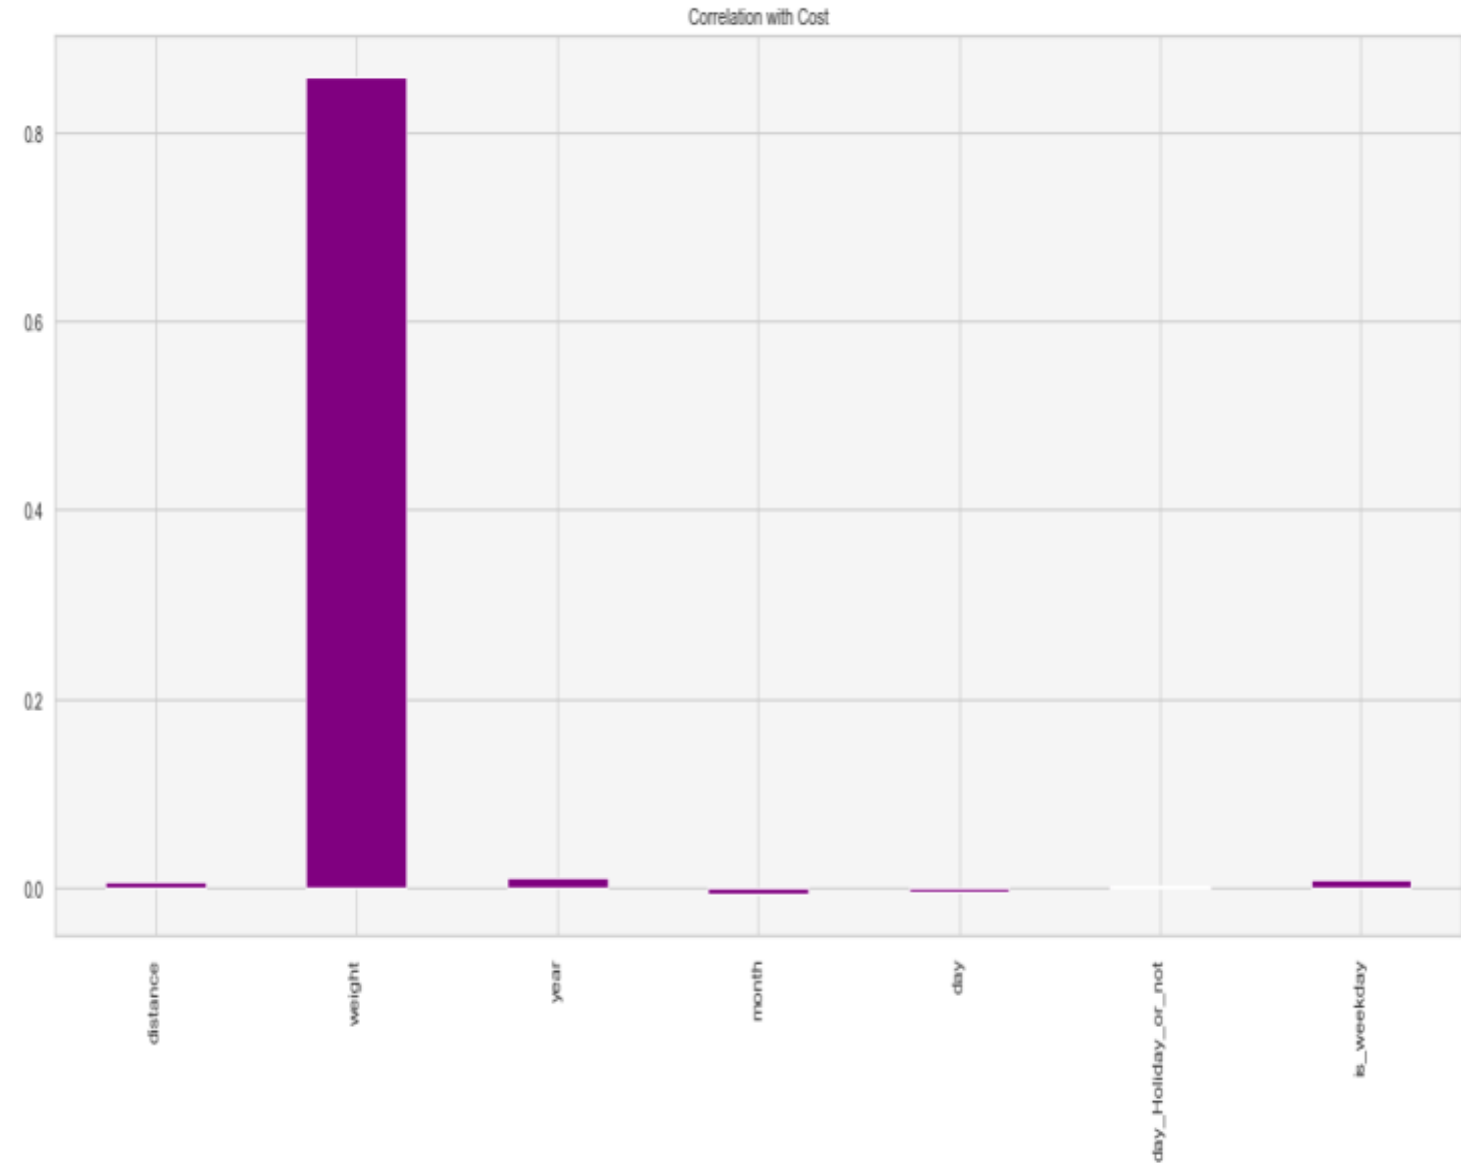
# Correlation heatmap

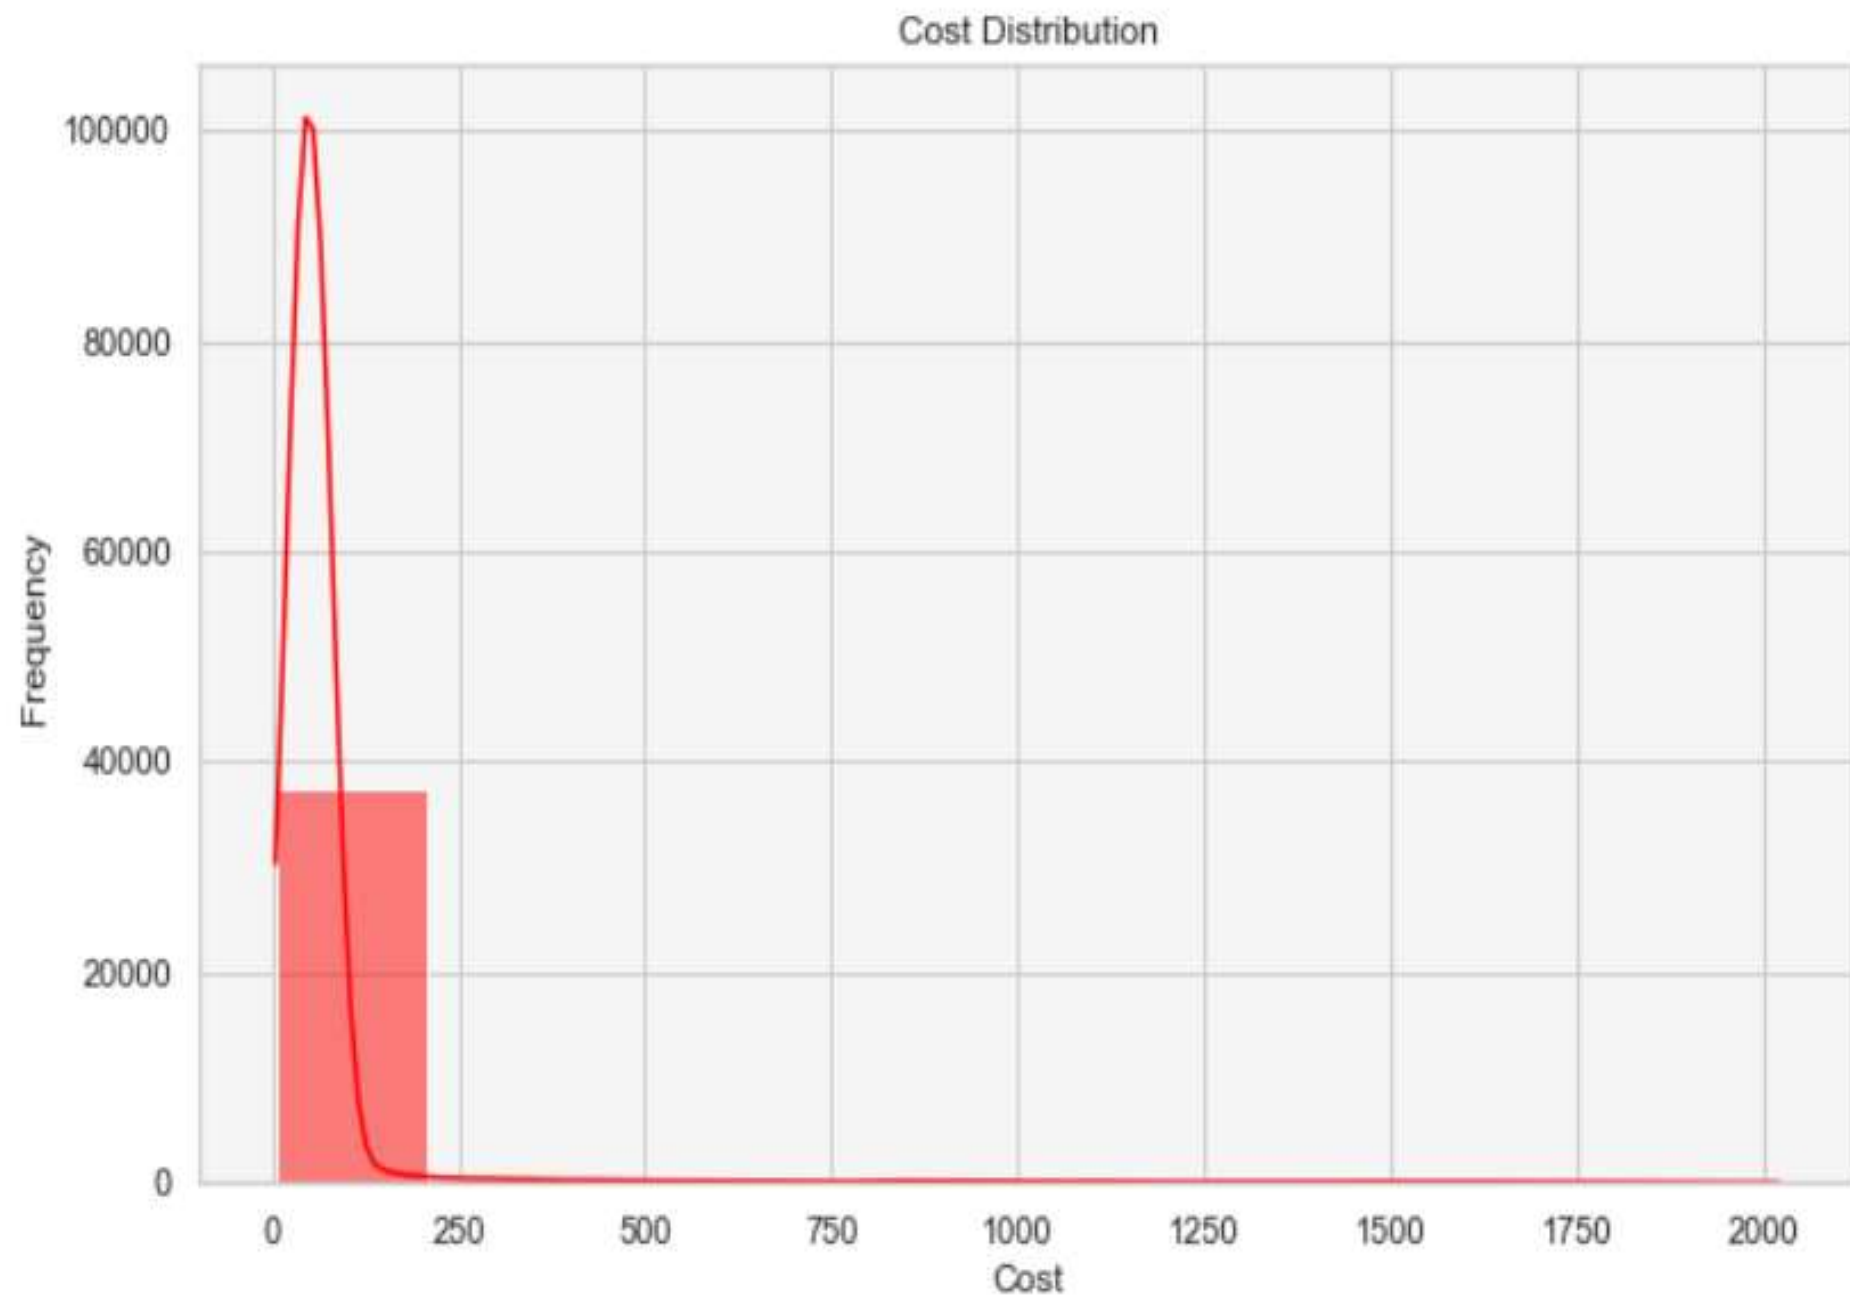


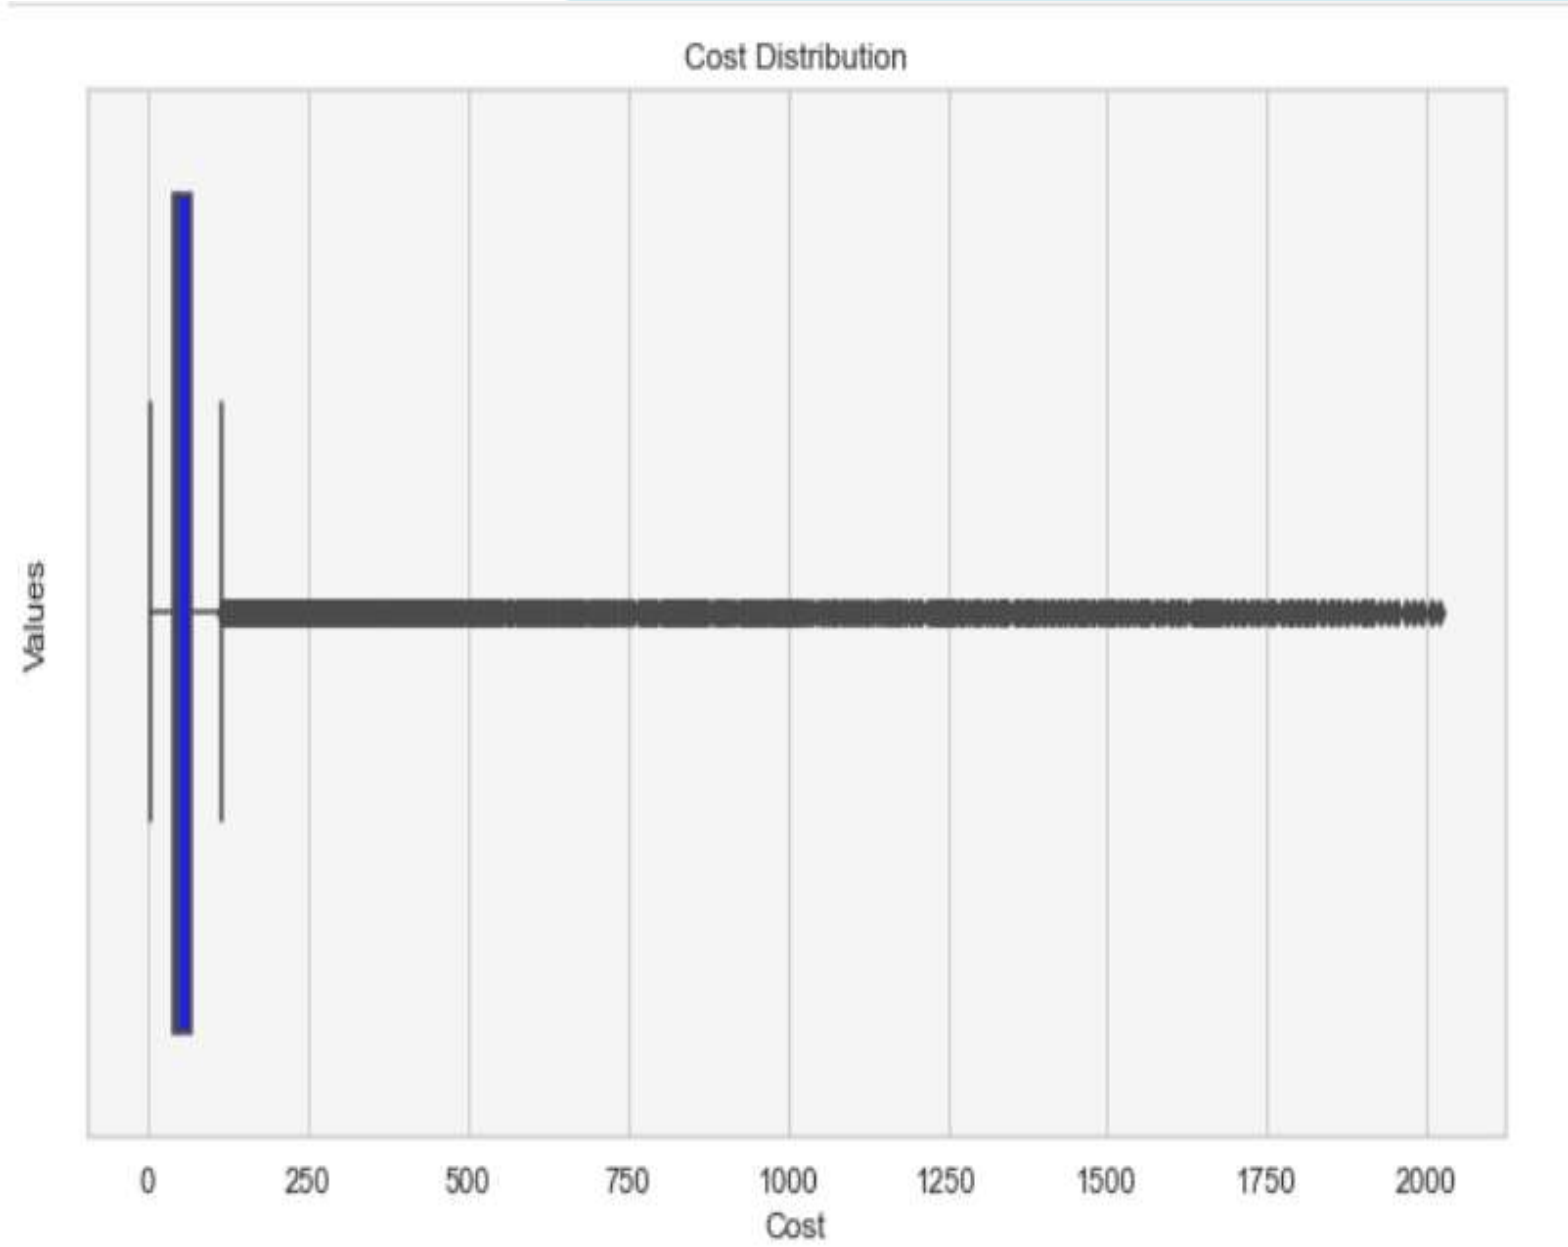


# Correlation plot with the target variable cost

Weight is highly correlated with the target variable



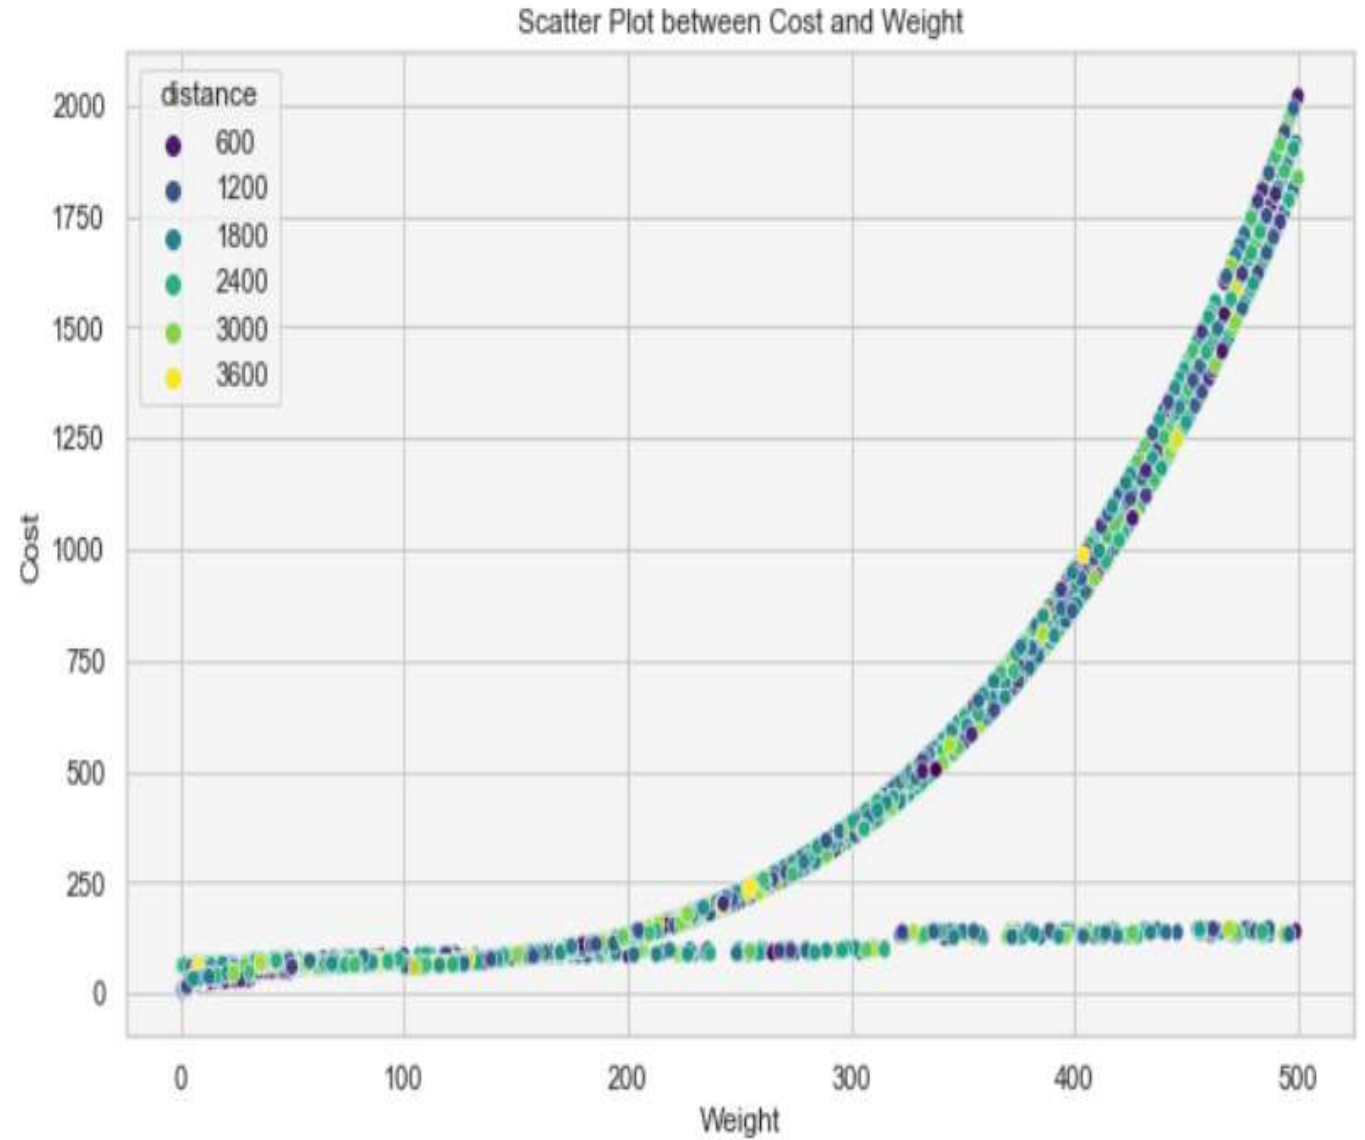


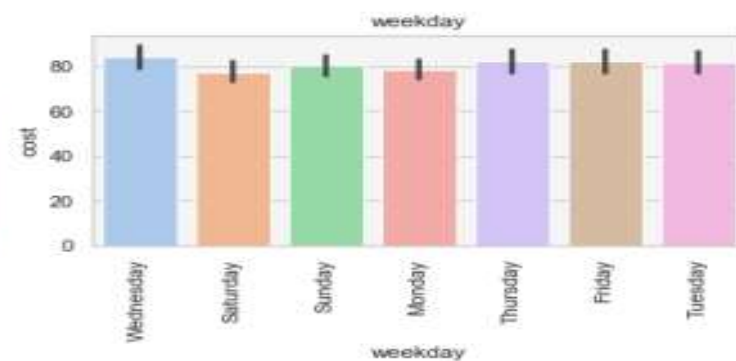
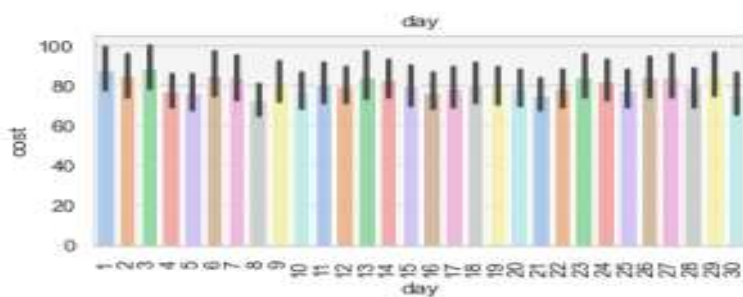
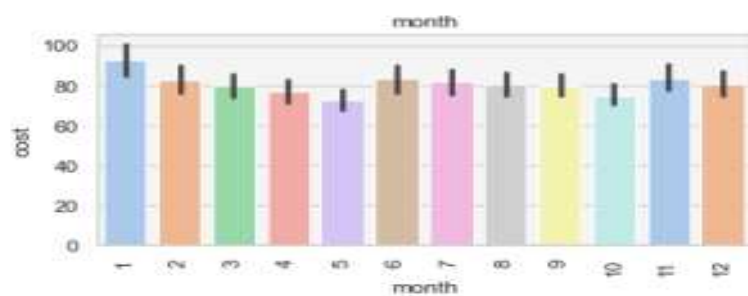
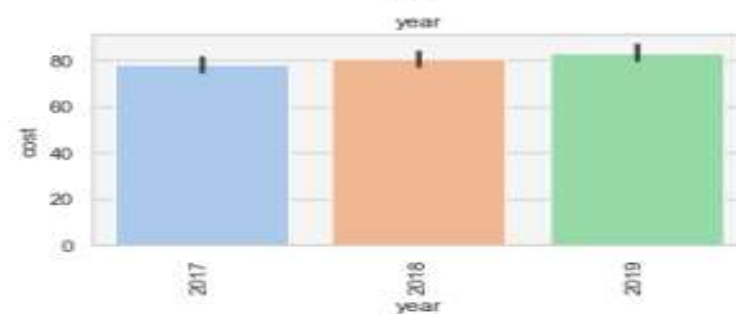
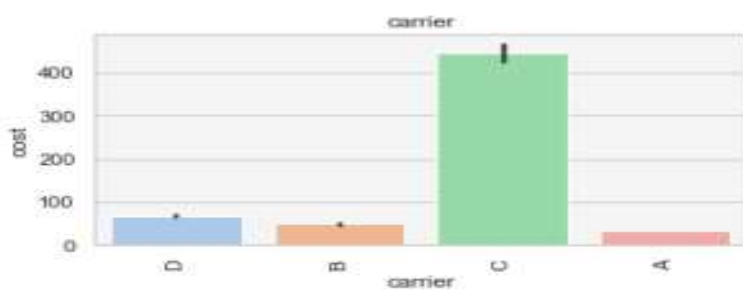
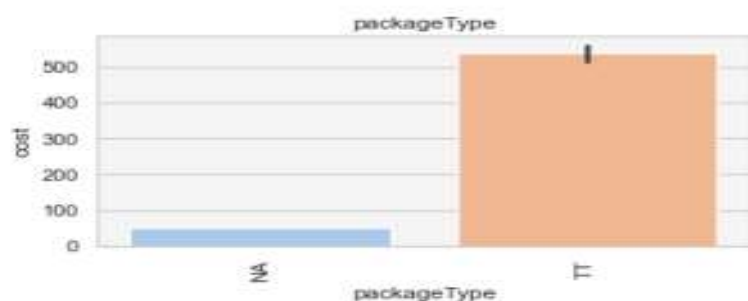
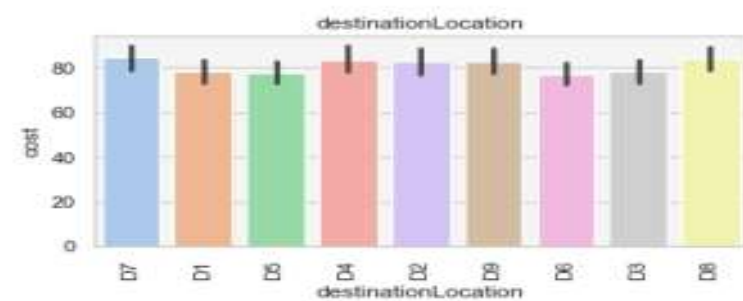
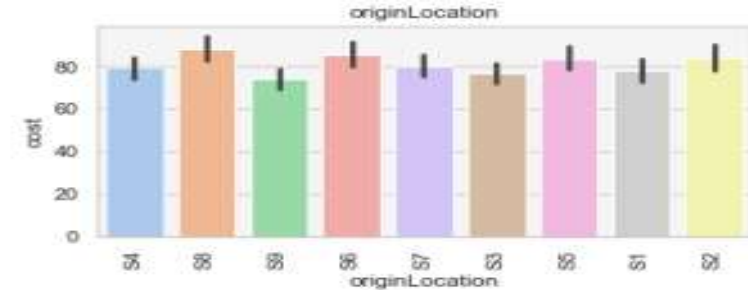
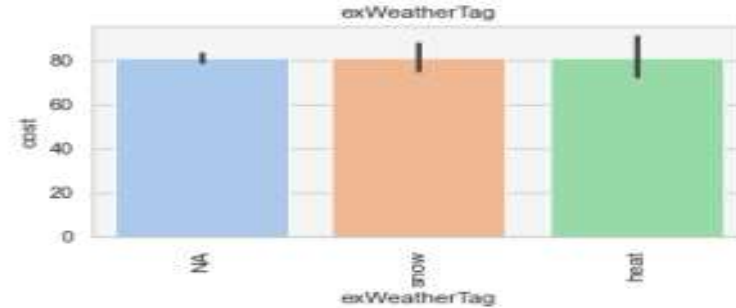
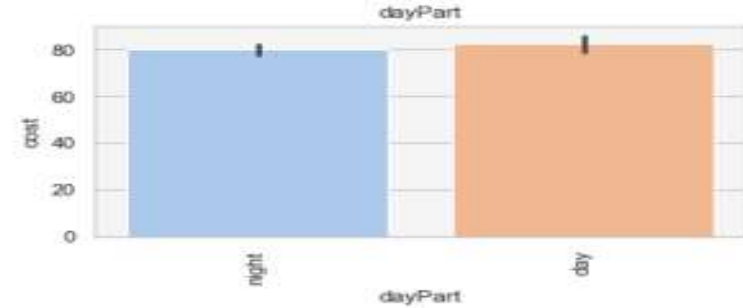




# Bivariate analysis

Performed on cost on weight variables





# Observations

LOCATIONS AND AGE ARE EQUALLY DISTRIBUTED.

AS WEIGHT INCREASES THE COST INCREASES.

MOST DELIVERIES ARE DONE DURING THE NIGHT

. WEIGHT IMPACTS THE COST MOST AND IT HAS HIGH CORRELATION WITH COST.

COST DOESN'T CHANGE BASED ON WEATHER AND DAYPART CONDITIONS.

COST IS HIGH IF THE PACKAGETYPE IS TT OR IF ITS CARRIED BY CARRIER C.

EXPEDITED PACKAGES HAVE A SLIGHTLY LOWER COST.

COST SLIGHTLY INCREASES BY YEAR.

AS THE YEAR STARTS THE COST IS HIGH AND IT DECREASES TILL MONTH 5 AND IT PICKS UP AFTER 6TH MONTH. OCTOBER HAS A SLIGHT DIP BUT IT PICKS UP AFTER THAT.



# *Modelling*



# Modelling

Feature encoding: Let's encode one hot encoding of the train, test and val together and later split them.

```
combined_df = pd.concat([X_train, X_test, test], axis=0)

# Perform one-hot encoding on the combined DataFrame
combined_df_encoded = pd.get_dummies(combined_df, columns=train.select_dtypes(exclude="number").columns.to_list())

# Split the combined DataFrame back into train, test, and val
train_encoded = combined_df_encoded.iloc[:len(X_train)]
test_encoded = combined_df_encoded.iloc[len(X_train):len(X_train) + len(X_test)]
val_encoded = combined_df_encoded.iloc[len(X_train) + len(X_test):]
```

# Training the dataset

- Linear Regression
- Ridge Regression
- XG Boost
- Cat Boost

```
# The below function trains the dataset for Linear Regression, Ridge regression, Xgboost and Catboost models and
def train_models(X_train, X_test, y_train, y_test):
    model_dict = {
        "linear": LinearRegression(),
        "Ridge": Ridge(alpha=0.2),
        ##"KNN": KNeighborsRegressor(n_jobs=-1, n_neighbors=4),
        "XGB": XGBRegressor(random_state=42),
        ## "light": LGBMRegressor(random_state=42),
        "Cat": CatBoostRegressor(random_state=42, loss_function='RMSE', verbose=False)
    }
    list1=[]
    dict1={}

    for model_name, model in model_dict.items():
        model.fit(X_train, y_train)
        pred = model.predict(X_test)
        num_predictors = X_train.shape[1]

        n = X_train.shape[0]

        adjusted_r_squared = 1 - (1 - r2_score(y_test, pred)) * (n - 1) / (n - num_predictors - 1)

        print(f"Training loss for model {model_name}   MSE : (mean_squared_error(y_test, pred, squared=False),)
        dict1[model_name]=[mean_squared_error(y_test, pred),sqrt(mean_squared_error(y_test, pred)),r2_score(y_t
                           mean_absolute_error(y_test, pred),model.score(X_test,y_test)*100,adjusted_r_squared]

    return dict1
```

# Training loss for the models

MODEL	MSE	RMSE	R-SQUARE	MAE	ACCURACY	ADJUSTED R SQUARE
linear	81.61	81.61	0.80	41.40	0.80	0.80
Ridge	81.61	81.61	0.80	41.40	0.80	0.80
XGB	1.71	1.71	0.99	0.34	0.99	0.99
Cat	1.58	1.58	0.99	0.39	0.99	0.99

```
]# Running the train_models function  
dict1=train_models(train_encoded,test_encoded, y_train, y_test)
```

```
Training loss for model linear  MSE : (81.6144663137134,),RMSE : (81.6144663137134,), R-Square : (0.8098722881  
943007,), MAE : (41.40249987247436,) , Accuracy : 0.8098722881943007 ,Adjusted R-Square : 0.8095583503881415  
Training loss for model Ridge  MSE : (81.61736842557802,),RMSE : (81.61736842557802,), R-Square : (0.809858766  
5302301,), MAE : (41.4006551439411,) , Accuracy : 0.8098587665302301 ,Adjusted R-Square : 0.8095448063971755  
Training loss for model XGB  MSE : (1.7149978499386376,),RMSE : (1.7149978499386376,), R-Square : (0.999916046  
5996152,), MAE : (0.3463397212158473,) , Accuracy : 0.9999160465996152 ,Adjusted R-Square : 0.9999159079762336  
  
Training loss for model Cat  MSE : (1.58649023994922,),RMSE : (1.58649023994922,), R-Square : (0.9999281567549  
223,), MAE : (0.39570604624943556,) , Accuracy : 0.9999281567549223 ,Adjusted R-Square : 0.999928038127761
```

# Initiating Catboost

- Since Catboost is our best model, we'll use gridsearch for hyperparameter tuning to find our best parameters.

```
In [104...  
# Instantiate CatBoostRegressor  
cbr = CatBoostRegressor()  
  
# Create a comprehensive grid with various hyperparameters  
grid = {  
    'depth': [6,9],  
    'learning_rate': [0.1],  
    'iterations': [100, 500],  
    'subsample': [0.5, 0.9],  
    'random_seed': [42]  
}  
  
# Define RMSE as the scoring metric  
rmse_scorer = make_scorer(lambda y_true, y_pred: np.sqrt(np.mean((y_true - y_pred) ** 2)), greater_is_better=False)  
  
# Instantiate GridSearchCV for CatBoostRegressor with the comprehensive grid and RMSE as the scoring metric  
gscv = GridSearchCV(estimator=cbr, param_grid=grid, scoring=rmse_scorer, cv=5)  
  
# Fit the grid search on the training data  
gscv.fit(train_encoded, y_train)  
  
# Print the best hyperparameters from the grid search  
print('Best Hyperparameters:', gscv.best_params_)  
  
# Use the best model from grid search for predictions  
best_model = gscv.best_estimator_  
y_pred = best_model.predict(test_encoded)  
  
# Calculate evaluation metrics on the test data  
rmse = sqrt(mean_squared_error(y_test, y_pred))  
mae = mean_absolute_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
# Print the evaluation metrics  
print('The RMSE score of the CatBoost is', rmse)  
print('The MAE score of the CatBoost is', mae)  
print('The R2 score of the CatBoost is', r2)  
print('The accuracy score of the CatBoost is', best_model.score(test_encoded, y_test) * 100)
```



```
482:   learn: 1.1024763   total: 2.11s   remaining: 74.3ms
483:   learn: 1.0994269   total: 2.11s   remaining: 69.9ms
484:   learn: 1.0986041   total: 2.12s   remaining: 65.5ms
485:   learn: 1.0949944   total: 2.12s   remaining: 61.1ms
486:   learn: 1.0947183   total: 2.13s   remaining: 56.8ms
487:   learn: 1.0929279   total: 2.13s   remaining: 52.4ms
488:   learn: 1.0909668   total: 2.13s   remaining: 48ms
489:   learn: 1.0884928   total: 2.14s   remaining: 43.6ms
490:   learn: 1.0873845   total: 2.14s   remaining: 39.3ms
491:   learn: 1.0862656   total: 2.15s   remaining: 34.9ms
492:   learn: 1.0852044   total: 2.15s   remaining: 30.5ms
493:   learn: 1.0834287   total: 2.15s   remaining: 26.2ms
494:   learn: 1.0809067   total: 2.16s   remaining: 21.8ms
495:   learn: 1.0789189   total: 2.16s   remaining: 17.5ms
496:   learn: 1.0748498   total: 2.17s   remaining: 13.1ms
497:   learn: 1.0736250   total: 2.17s   remaining: 8.72ms
498:   learn: 1.0726284   total: 2.18s   remaining: 4.36ms
499:   learn: 1.0720685   total: 2.18s   remaining: 0us
```

```
Best Hyperparameters: {'depth': 6, 'iterations': 500, 'learning_rate': 0.1, 'random_seed': 42, 'subsample': 0.5}
```

```
The RMSE score of the CatBoost is 1.8319448087674535
```

```
The MAE score of the CatBoost is 0.5125643984315946
```

```
The R2 score of the CatBoost is 0.9999042065306653
```

```
The accuracy score of the CatBoost is 99.99042065306652
```

```
In [105... print('Best Hyperparameters:', gscv.best_params_)
```

```
Best Hyperparameters: {'depth': 6, 'iterations': 500, 'learning_rate': 0.1, 'random_seed': 42, 'subsample': 0.5}
```

## Let's use adaboost to improve on the previous catboost.

```
8]: from sklearn.ensemble import AdaBoostRegressor

ada = AdaBoostRegressor(base_estimator = model,n_estimators=200,learning_rate=0.1)

ada.fit(train_encoded, y_train)
y_pred = ada.predict(test_encoded)

rmse = sqrt(mean_squared_error(y_test, y_pred))

mae = mean_absolute_error(y_test, y_pred)
r2=r2_score(y_test, y_pred)

print('The RMSE score of the catboost is ',rmse)
print('The MAE score of the catboost ',mae)
print('The R2 score of the catboost is ',r2)
print('The accuracy score of the catboost is ',ada.score(test_encoded,y_test)*100)
```

The RMSE score of the catboost is 0.6795211189527792

The MAE score of the catboost 0.1251449807349938

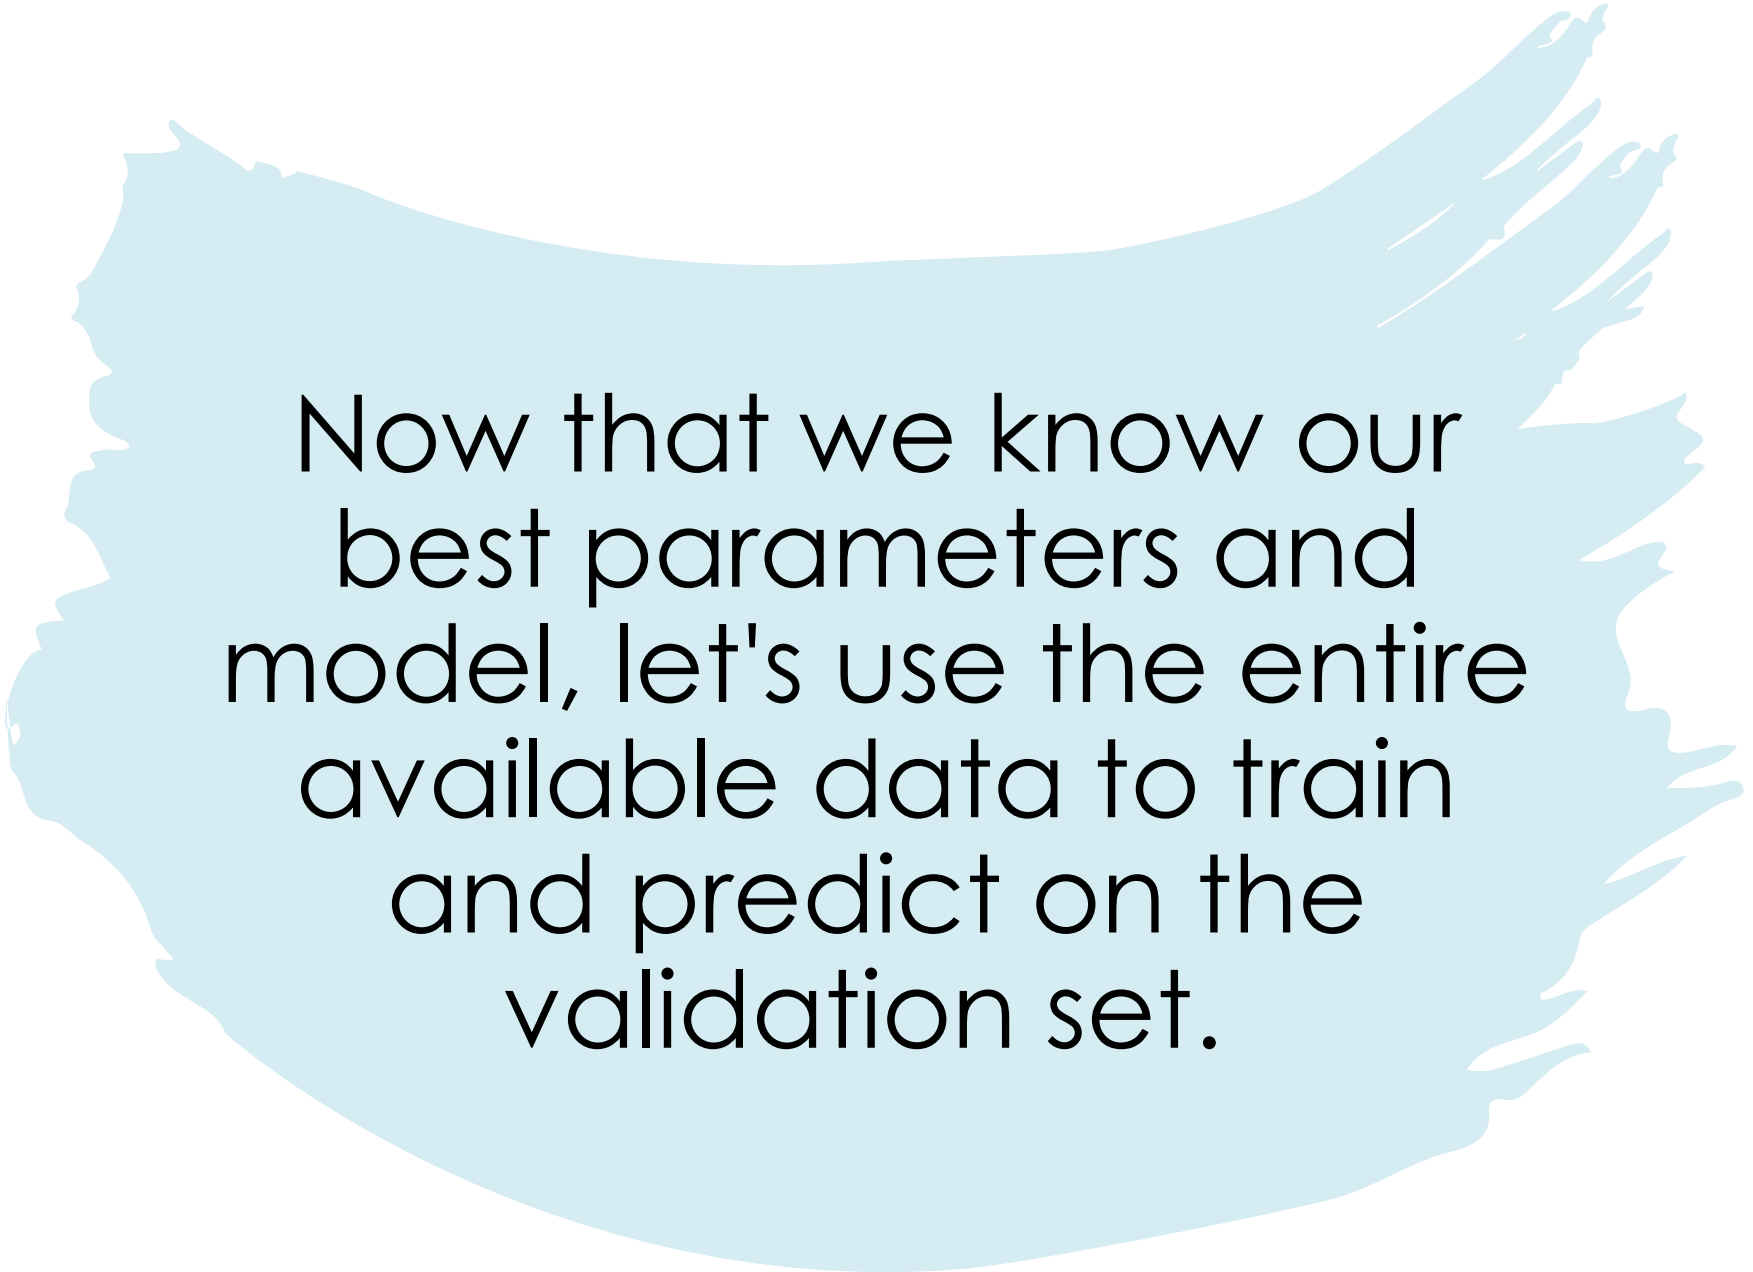
The R2 score of the catboost is 0.9999868199502696

The accuracy score of the catboost is 99.99868199502696

# Ada boost

This is the best model on our test set and had an rmse of 0.6 on val set also.

```
val_set=pd.DataFrame(val_trip)
val_set['cost']=ada.predict(val_encoded)
val_set.to_csv('Submission6.csv')
```

A light blue brushstroke background with a textured, painterly appearance, featuring various shades of blue and white strokes that create a sense of movement and depth.

Now that we know our  
best parameters and  
model, let's use the entire  
available data to train  
and predict on the  
validation set.



# Combining x\_train and x\_test

Ada improved the RMSE and after a submitting it to Kaggle we had the highest personal best of 0.4 public score and 0.3 private score

```
# Combining x_train and x_test
combined_x_train = pd.concat([train_encoded, test_encoded], axis=0)
combined_y_train = pd.concat([y_train, y_test], axis=0)

ada.fit(combined_x_train, combined_y_train)
y_pred = ada.predict(test_encoded)

rmse = sqrt(mean_squared_error(y_test, y_pred))

mae = mean_absolute_error(y_test, y_pred)
r2=r2_score(y_test, y_pred)

print('The RMSE score of the catboost is ',rmse)
print('The MAE score of the catboost ',mae)
print('The R2 score of the catboost is ',r2)
print('The accuracy score of the catboost is ',ada.score(test_encoded,y_test)*100)
```

```
The RMSE score of the catboost is  0.0647618394141029
The MAE score of the catboost  0.037675440082020006
The R2 score of the catboost is  0.9999998802847485
The accuracy score of the catboost is  99.99998802847485
```

```
val_set=pd.DataFrame(val_trip)
val_set['cost']=ada.predict(val_encoded)
val_set.to_csv('Submission7.csv')
```

# ANN using Tensorflow for Regression

LET'S USE A NEURAL NETWORK WITH THREE  
LAYERS AND RELU AS THE ACTIVATION  
FUNCTION

```
In [45]: ## Let's use a neural network with three layers and relu as the activation function
nn_model = tf.keras.Sequential([
    tf.keras.layers.Dense(units = 64, activation = tf.nn.relu, input_shape = [train_encoded.shape[1]]),
    tf.keras.layers.Dense(units = 64, activation = tf.nn.relu),
    tf.keras.layers.Dense(units = 1)
])
```

```
In [51]: nn_model.compile(loss = 'mse', optimizer = tf.keras.optimizers.RMSprop(0.001), metrics = ['mae', 'mse'])
```

```
In [52]: nn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 64)	2944
-----		
dense_1 (Dense)	(None, 64)	4160
-----		
dense_2 (Dense)	(None, 1)	65
=====		

Total params: 7,169

Trainable params: 7,169

Non-trainable params: 0

```
In [53]: history = nn_model.fit(train_encoded, y_train, epochs = 500, validation_data=(test_encoded, y_test))
```

```

Epoch 495/500
854/854 [=====] - 4s 4ms/step - loss: 5.1839 - mae: 0.9507 - mse: 5.1839 - val_loss: 13.0082 - val_mae: 1.1555 - val_mse: 13.0082
Epoch 496/500
854/854 [=====] - 4s 5ms/step - loss: 5.2333 - mae: 0.9378 - mse: 5.2333 - val_loss: 7.2455 - val_mae: 1.4560 - val_mse: 7.2455
Epoch 497/500
854/854 [=====] - 4s 4ms/step - loss: 5.2347 - mae: 0.9386 - mse: 5.2347 - val_loss: 6.7140 - val_mae: 1.1064 - val_mse: 6.7140
Epoch 498/500
854/854 [=====] - 4s 4ms/step - loss: 5.1578 - mae: 0.9328 - mse: 5.1578 - val_loss: 4.4918 - val_mae: 0.8519 - val_mse: 4.4918
Epoch 499/500
854/854 [=====] - 4s 4ms/step - loss: 5.2732 - mae: 0.9337 - mse: 5.2732 - val_loss: 10.3131 - val_mae: 1.2639 - val_mse: 10.3131
Epoch 500/500
854/854 [=====] - 4s 4ms/step - loss: 4.9170 - mae: 0.9308 - mse: 4.9170 - val_loss: 5.1875 - val_mae: 0.8679 - val_mse: 5.1875

```

```

In [54]: hist = pd.DataFrame(history.history)
hist.tail()

```

```

Out[54]:

```

	loss	mae	mse	val_loss	val_mae	val_mse
<b>495</b>	5.233289	0.937784	5.233289	7.245465	1.456009	7.245465
<b>496</b>	5.234665	0.938590	5.234665	6.713998	1.106439	6.713998
<b>497</b>	5.157828	0.932813	5.157828	4.491841	0.851876	4.491841
<b>498</b>	5.273247	0.933667	5.273247	10.313062	1.263892	10.313062
<b>499</b>	4.916996	0.930770	4.916996	5.187454	0.867891	5.187454

## Printing the scores of Neural Network results

The RMSE score is 2.27759863217814

The MAE score 0.8678907700156109

The R2 score is 0.9998519305313268

```
## Printing the NN results
y_pred = nn_model.predict(test_encoded)
rmse = sqrt(mean_squared_error(y_test, y_pred))

mae = mean_absolute_error(y_test, y_pred)
r2=r2_score(y_test, y_pred)

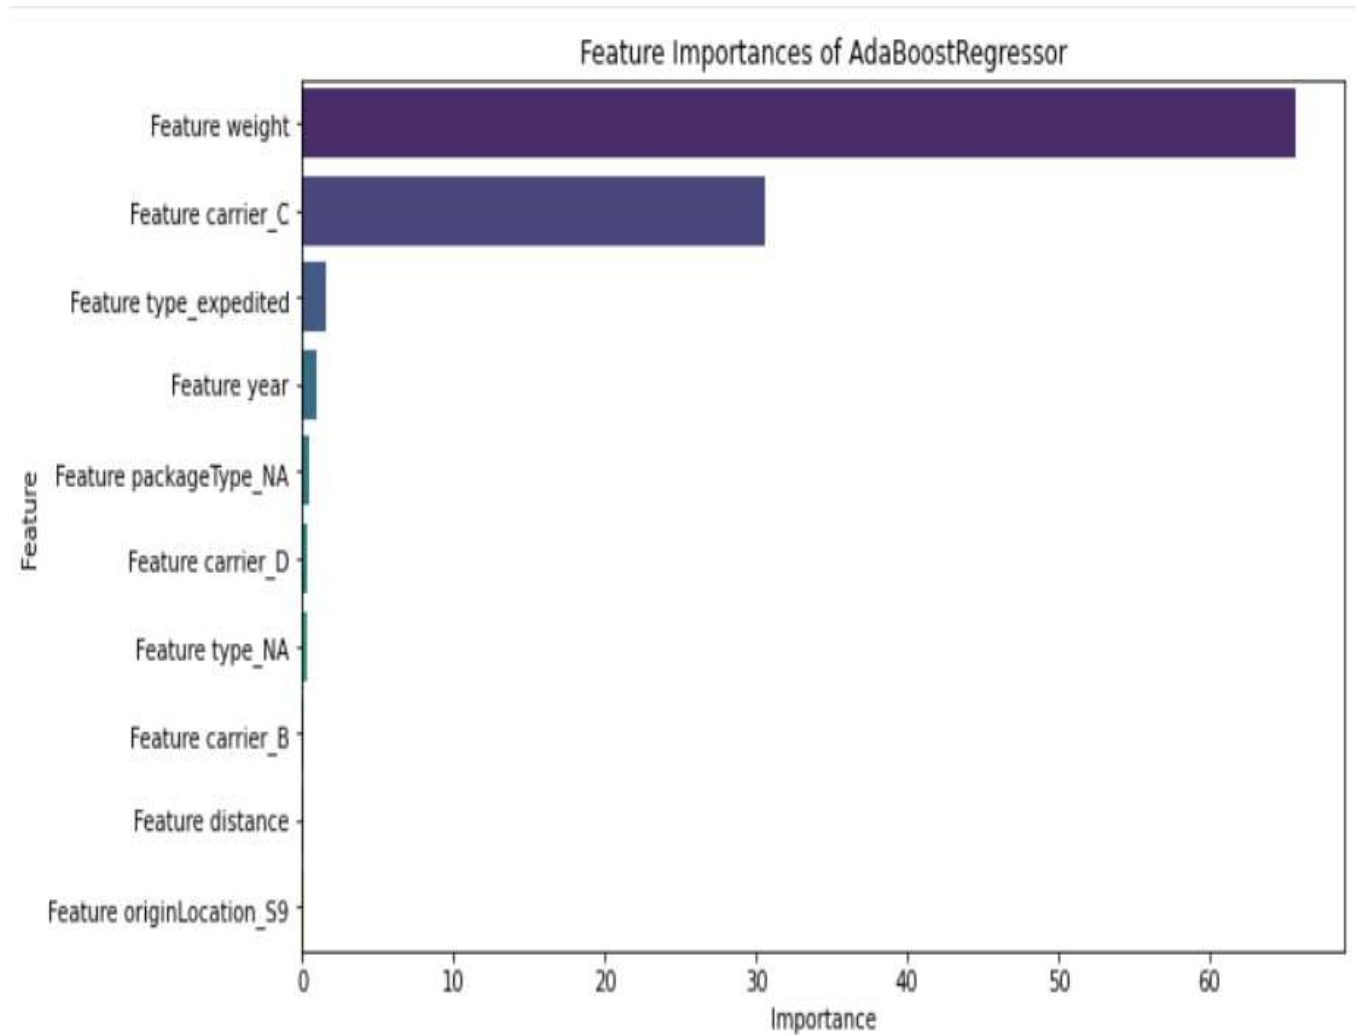
print('The RMSE score of the neural network is ',rmse)
print('The MAE score of the neural network ',mae)
print('The R2 score of the neural network is ',r2)
```

```
The RMSE score of the neural network is  2.27759863217814
The MAE score of the neural network  0.8678907700156109
The R2 score of the neural network is  0.9998519305313268
```



# Feature Importance

- As we found from the EDA weight and Carrier C and expedited are the best predictors of the cost.






# Conclusion

- We chose catboost and adaboost as our final models, we used base models such as Linear regression, Ridge Regression, Xgboost and Catboost.
- The model was evaluated based on RMSE, MAE, R-Square and adjusted R-square.



# Score in Kaggle

- Our competition kaggle score was 1.2164 and our late submitted score improved to 0.4 RMSE on the validation set

Submission and Description		Private Score ⓘ	Public Score ⓘ	Selected
 <b>Submission7.csv</b> Complete (after deadline) · KapileshAp · now		<b>0.4282</b>	<b>0.3472</b>	<input type="checkbox"/>
 <b>Submission6.csv</b> Complete (after deadline) · KapileshAp · 1h ago		<b>0.6502</b>	<b>0.49565</b>	<input type="checkbox"/>
 <b>Submission6.csv</b> Complete · KapileshAp · 9h ago		<b>1.22271</b>	<b>1.12164</b>	<input checked="" type="checkbox"/>

# References

- <https://www.kaggle.com/competitions/cost-prediction-for-logistic-company-fall2023/overview>





*ANY  
QUESTIONS?*





Thank you